

Further Reading

You Belong in the Loop — Sam Seim (they/them) — 2027 Nebraska Women in STEM Conference

Every citation on this page was checked against the publisher's record before it was printed, which is a low bar and one that reading lists in this field clear less often than you would like.

Where I could not verify something, it is not here. There is one such case and I have written up why, at the bottom, because the way it failed is the most useful thing on the page.

THE FRAMEWORK BEHIND THE SESSION

Olmanson, J., Mwayaona, J., Seim, S., & Stevens, K. (2026). *Learning Technology Design and Prototyping in the Age of Agentic AI: The 10(D) 8(Re) Framework*. EdMedia + Innovate Learning 2026, Edinburgh, Scotland. AACE.

The design framework the dissertation instrument is situated within. I am a co-author on this paper.

ON THE MECHANISM — WHY SYSTEMS ARE WORSE AT THE TAIL

Koenecke, A., Nam, A., Lake, E., Nudell, J., Quartey, M., Mengesha, Z., Toups, C., Rickford, J. R., Jurafsky, D., & Goel, S. (2020). Racial disparities in automated speech recognition. PNAS, 117(14), 7684–7689. doi.org/10.1073/pnas.1915768117

Start here. Five commercial speech recognizers, measurably worse for Black speakers than white speakers on matched material. If you take one paper from this page, take this one — it is the accommodation-by-averaging argument in its most literal form, in the exact modality I depend on.

Buolamwini, J., & Gebru, T. (2018). Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. PMLR 81, 77–91. proceedings.mlr.press/v81/buolamwini18a.html

The study that established the method: disaggregate by group and the average dissolves into something much less flattering.

Raji, I. D., & Buolamwini, J. (2019). Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products. AIES '19, 429–435.

What happened after Gender Shades — the audit as an instrument of change rather than a report. There is a 2023 follow-up in CACM; they are different papers.

Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? FAccT '21, 610–623.

The scale argument. Read it for the section on who is and is not represented in a training corpus, which is the whole of the middle-of-the-distribution claim.

ON DISABILITY SPECIFICALLY

Bennett, C. L., & Keyes, O. (2020). What is the point of fairness? Disability, AI and the complexity of justice. ACM SIGACCESS Accessibility and Computing, 125. Preprint: arXiv:1908.01024

The paper that refuses the easy version. Disability is not a category you can balance a dataset across — it is heterogeneous by definition, which breaks most fairness methods at the root rather than at the margin.

Whittaker, M., Alper, M., Bennett, C. L., Hendren, S., Kaziunas, L., Mills, M., Morris, M. R., Rankin, J., Rogers, E., Salas, M., & West, S. M. (2019). Disability, Bias, and AI. AI Now Institute, New York University. ainowinstitute.org

A report, not a paper, and the most useful single overview I know for handing to a colleague who is new to the argument.

Costanza-Chock, S. (2020). Design Justice: Community-Led Practices to Build the Worlds We Need. MIT Press. Open access: doi.org/10.7551/mitpress/12255.001.0001

The book that reframes the question from what does this cost to who is in the room when it is designed. Free, legally, from the publisher.

ON EVALUATION PRACTICE — THE PART YOU CAN ACTUALLY ADOPT MONDAY

Ribeiro, M. T., Wu, T., Guestrin, C., & Singh, S. (2020). Beyond Accuracy: Behavioral Testing of NLP Models with CheckList. ACL 2020, 4902–4912.

The wrong-answer file with a method attached. Best paper at ACL that year and the most directly practical thing here.

Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model Cards for Model Reporting. FAT* '19, 220–229.

What a vendor should be able to tell you, in a format you can ask for by name.

Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., & Crawford, K. (2021). Datasheets for datasets. Communications of the ACM, 64(12), 86–92.

The companion question about the data rather than the model.

Suresh, H., & Gutttag, J. V. (2021). A Framework for Understanding Sources of Harm throughout the Machine Learning Life Cycle. EAAMO '21.

Where harm enters, stage by stage. Useful for locating your own build on the map rather than worrying about all of it at once.

Shelby, R., Rismani, S., Henne, K., Moon, A., Rostamzadeh, N., Nicholas, P., Yilla-Akbari, N., Gallegos, J., Smart, A., Garcia, E., & Virk, G. (2023). Sociotechnical Harms of Algorithmic Systems: Scoping a Taxonomy for Harm Reduction. AIES '23, 723–741.

A taxonomy, which sounds dry and is the thing that lets you write a specific risk into a specific document.

STANDARDS YOU CAN PUT IN A MEMO

NIST (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1. nist.gov/itl/ai-risk-management-framework

Voluntary, government-issued, and vocabulary your risk function already accepts. The Generative AI Profile (NIST AI 600-1, 2024) is the companion.

W3C (2024). Web Content Accessibility Guidelines (WCAG) 2.2. W3C Recommendation, 12 December 2024. [w3.org/TR/WCAG22/](https://www.w3.org/TR/WCAG22/)

2.2 is the current standard. WCAG 3.0 is a Working Draft and is not — a fair number of confident blog posts say otherwise, and they are wrong.

A NUMBER I WENT LOOKING FOR AND AM NOT GOING TO GIVE YOU

There is a figure that circulates constantly: retrofitting accessibility costs ten times what building it in would have. It is quoted in vendor blogs, agency decks, and procurement business cases. I wanted it for this session.

It does not have a source. Traced back, it is the general software cost-of-change curve — Boehm, 1981, drawn from 1970s waterfall projects and contested since — with the word “accessibility” attached to it somewhere along the way. There is no accessibility-specific study behind the multiplier. The nearest real document is a 2005 World Bank education note about the cost of accessible school buildings, which is qualitative, is about physical construction, and does not cite a source for its own figure. W3C’s own business case for digital accessibility pointedly declines to make the comparison.

So the number is off this page, and it is off the slides. Which is the session’s own rule applied to the session: quantify only in the unit the system actually keeps, and if the real unit is not known, find it or do not count. A fabricated metric is not a metric — it is a liability with a decimal point, and it will be the thing you get asked about.

The defensible version of the claim survives without it: designing with disabled users from the start changes what gets built, not merely what it costs. That one you can stand behind in a room, and Costanza-Chock above does the work.

For the record, I do have a number. It is just mine rather than the world’s. I fought my own dissertation tracker for about a semester and a half trying to retrofit voice into something that was never built for it, then scrapped it, wrote down what it actually had to do, and rebuilt it over a weekend. A semester and a half, and then a weekend. That is not a statistic and I would not put it in anybody’s business case — it is one person’s kept unit, reported as one person’s kept unit, which is the entire distinction this page is about.

Nothing on this page names a product. The products will change; these questions have not.

SAM SEIM — THEY/THEM

sam@drseim.com

www.drseim.com