

The Workflow Audit

You Belong in the Loop — Sam Seim (they/them) — 2027 Nebraska Women in STEM Conference

THE RULE

Delegate the task whose output you can verify faster than you could produce it.

If verification costs more than doing it yourself, delegation is a loss you haven't priced yet. That doesn't mean don't — it means book the check.

STEP 1 — THREE MINUTES, ALONE, IN SILENCE

Pick one real task you did last week. Not a hypothetical, not the task you wish you had. Fill in the two lines below before you touch the table — writing down what wrong looks like, first, is the whole trick and the step everybody skips.

The task

What a wrong result looks like

STEP — WHAT I ACTUALLY DO	WHERE THE JUDGMENT IS	MINS TO DO	MINS TO CHECK	VERDICT
Reconcile the monthly source file against the ledger	Deciding what a mismatch means	90	10 — the totals either tie or they don't	Delegate the matching. Keep the mismatch call.

STEP 2 — FIVE MINUTES, IN PAIRS

Say the numbers out loud. For each step, ask your neighbor: could you verify a machine's version of this faster than you could do it yourself? Make them give you a number in minutes. A shrug is not a number.

If you get stuck: what would you have to look at to know it was right? Now, how long does looking at that actually take? That is the number.

STEP 3 — THREE MINUTES, OUT LOUD

One thing. Not your whole task. One step that turned out to be genuinely delegable, or — better — one that looked delegable right up until you priced the check.

Four people. If you don't get called on, the row is still yours to act on Monday.

The Reference Card

Everything worth keeping from this session. Nothing here names a product, because the product will change.

PROMPT, WORKFLOW, AGENT — THREE POSITIONS ON ONE TRADE

	WHAT IT COSTS	HOW IT FAILS	RIGHT WHEN
Prompt	Your attention, per use, forever.	Immediately, in front of you.	The task is rare, or the judgment is the actual product.
Workflow	Built once, maintained forever.	Loudly. It breaks and it stops.	The task repeats and the steps hold still. Most people need this one.
Agent	The build, plus verification that never ends.	Quietly, plausibly, and at volume.	The steps genuinely can't be known in advance — and you can afford the check.

ONE QUESTION, TWO LEDGERS

Ask them across, not down. They are the same four questions.

- Where does it fail? • **Who is it worse for?**
How would I know? • **How would I know?**
What does the check cost? • **What does that check cost?**
Who signs off? • **Who pays when nobody checks?**

The left column has a budget line. The right column has a volunteer.

FOUR PATTERNS FOR A LOCKED-DOWN SHOP

Move the reasoning, not the data. Send the schema, the column names, a synthetic row that resembles no real person. What comes back is a transform you run inside the wall.

Automate the boring eighty. Most of what looks like an AI problem is a parsing problem wearing a hat. Reserve the model for the part that needs judgment.

Get the no in writing. Ask which specific control you're tripping. "No AI" is rarely the policy; it's the summary of one, and the written version has edges you can build to.

Log it like you'll be asked. Input, version, model, timestamp, who approved. If you can't reconstruct a decision in two years, you don't have a system, you have a habit.

IN THE LOOP IS NOT AUTHORITY TO STOP IT

Real oversight needs all three. Two out of three is decoration.

Information. You can see what the system did, not only what it produced. A summary of an action is not a record of it.

Standing. Stopping it is a normal action, not an escalation with your name on it. If halting costs you politically, you are decoration.

Time. The check is on the schedule with a name against it. If review is unpaid overtime, there is no review — there's a checkbox.

SIX EVALUATION HABITS

- 1 Write the check before the build.** If you can't say what wrong looks like, you are not ready to automate.
- 2 Keep a wrong-answer file.** Every failure, verbatim, with the input that caused it.
- 3 Run one input through two models.** Divergence is a free map of where the task is underspecified.
- 4 Sample by group, not at random.** Random sampling finds the average failure. That is not the one that ends up in a complaint.
- 5 Version the prompt like code.** It is code. It has a blast radius and you changed it on a Friday.
- 6 Re-run the old failures every time.** That's a regression test. You already know how to do this.

SELLING IT UPWARD

"We can't verify this output" and "this output is worse for a third of our customers" are the same memo. One of them gets funded — write that one. The risk framing isn't a trick; it is the true description that happens to have a budget code.